

ACT_4

October 21, 2023

- 1 Team Number: 7
- 2 Team Captain: Jacob Silva
- 3 Team Members: Joel Hurtado
- 4 Juliana Steele
- 5 1. Read the data into your software system

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import norm
import statsmodels.api as sm
import pylab as py
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
from scipy import stats
```

```
[ ]: data = pd.read_csv("HeartDisease.csv")

# chd is target variable
data.head()
```

```
[ ]: 
```

	names	sbp	tobacco	ldl	adiposity	famhist	typea	obesity	alcohol	\
0	1	160	12.00	5.73	23.11	Present	49	25.30	97.20	
1	2	144	0.01	4.41	28.61	Absent	55	28.87	2.06	
2	3	118	0.08	3.48	32.28	Present	52	29.14	3.81	
3	4	170	7.50	6.41	38.03	Present	51	31.99	24.26	
4	5	134	13.60	3.50	27.78	Present	60	25.99	57.34	

	age	chd
0	52	1
1	63	1
2	46	0
3	58	1

6 2. Examine univariate statistics for the following variables: sbp, tobacco, ldl, adiposity, typea, obesity, alcohol, and age

```
[ ]: df = data[['sbp', 'tobacco', 'ldl', 'adiposity', 'typea', 'obesity', 'alcohol', 'age']]

stats = df.describe()
stats = stats.transpose()

print(stats)
```

	count	mean	std	min	25%	50%	75%	\
sbp	462.0	138.326840	20.496317	101.00	124.0000	134.000	148.0000	
tobacco	462.0	3.635649	4.593024	0.00	0.0525	2.000	5.5000	
ldl	462.0	4.740325	2.070909	0.98	3.2825	4.340	5.7900	
adiposity	462.0	25.406732	7.780699	6.74	19.7750	26.115	31.2275	
typea	462.0	53.103896	9.817534	13.00	47.0000	53.000	60.0000	
obesity	462.0	26.044113	4.213680	14.70	22.9850	25.805	28.4975	
alcohol	462.0	17.044394	24.481059	0.00	0.5100	7.510	23.8925	
age	462.0	42.816017	14.608956	15.00	31.0000	45.000	55.0000	

	max
sbp	218.00
tobacco	31.20
ldl	15.33
adiposity	42.49
typea	78.00
obesity	46.58
alcohol	147.19
age	64.00

7 3. Produce histogram of each of the following variables with imposing normal curve: sbp, tobacco, ldl, adiposity, typea, obesity, alcohol, and age.

```
[ ]: axes = df.hist(grid=False, figsize=(10, 8), layout=(3, 3))


axes
```

```
[ ]: array([[<Axes: title={'center': 'sbp'}>,
          <Axes: title={'center': 'tobacco'}>,
          <Axes: title={'center': 'ldl'}>],
          [<Axes: title={'center': 'adiposity'}>,
          <Axes: title={'center': 'typea'}>],
          []])
```

```

<Axes: title={'center': 'obesity'}>],
[<Axes: title={'center': 'alcohol'}>,
<Axes: title={'center': 'age'}>, <Axes: >]], dtype=object)

```



```

[ ]: num_cols = len(df.columns)

fig, axes = plt.subplots(3, 3, figsize=(10, 8))

axes = axes.ravel()

for i in range(min(num_cols, 9)):
    column_data = df.iloc[:, i]

    axes[i].hist(column_data, bins=10, density=True, alpha=0.6, color='b')

    mu, std = column_data.mean(), column_data.std()
    xmin, xmax = axes[i].get_xlim()
    x = np.linspace(xmin, xmax, 100)

```

```

p = norm.pdf(x, mu, std)

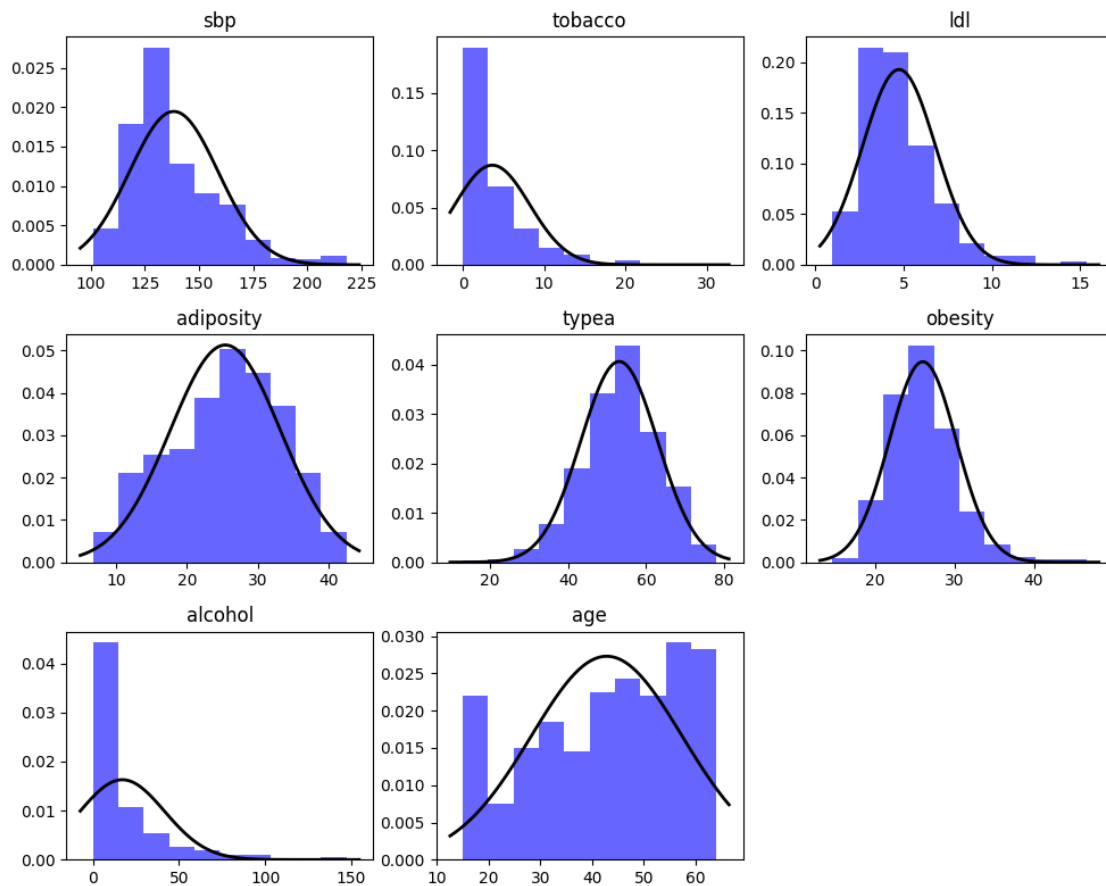
axes[i].plot(x, p, 'k', linewidth=2)

axes[i].set_title(df.columns[i])

for i in range(num_cols, 9):
    axes[i].axis('off')

plt.tight_layout()
plt.show()

```

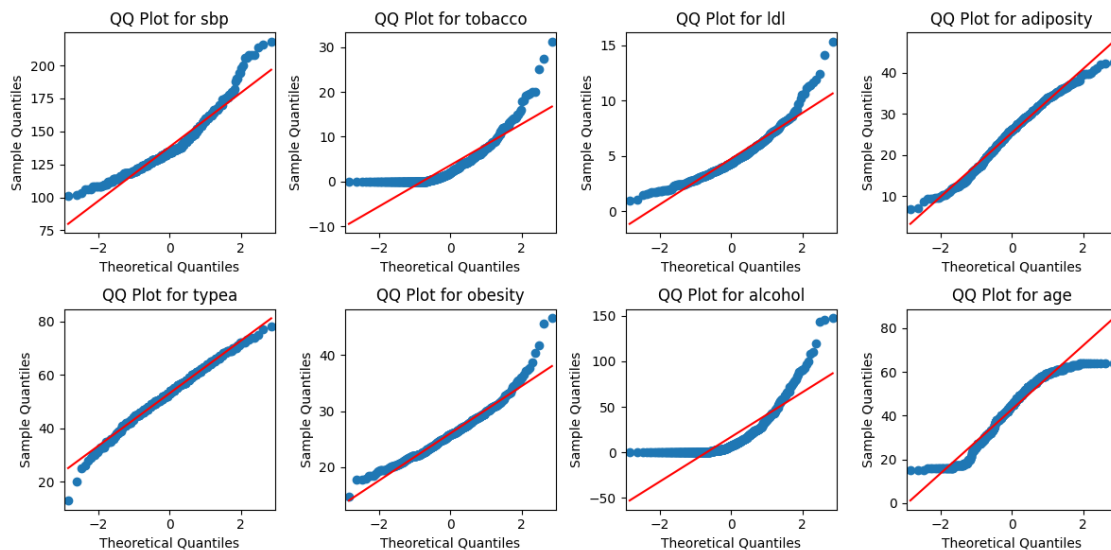


8 4. Produce quantile plot of each of the following variables: sbp, tobacco, ldl, adiposity, typea, obesity, alcohol, and age

```
[ ]: fig, axes = plt.subplots(2, 4, figsize=(12, 6))
axes = axes.ravel()

for i, column in enumerate(df.columns):
    ax = axes[i]
    sm.qqplot(df[column], line='s', ax=ax)
    ax.set_title(f'QQ Plot for {column}')

plt.tight_layout()
plt.show()
```



9 5. Build a logistic regression model with all predictors

```
[ ]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report
X = df
y = data['chd']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

# in LogisticRegression x=10.0 changes the regularization strength
```

```

model = LogisticRegression(solver='liblinear', random_state=0).fit(X,y)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

confusion_matrix(y_test, y_pred)

```

```
[ ]: array([[49, 10],
          [19, 15]])
```

```
[ ]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.72	0.83	0.77	59
1	0.60	0.44	0.51	34
accuracy			0.69	93
macro avg	0.66	0.64	0.64	93
weighted avg	0.68	0.69	0.68	93

10 6. Perform power transformation on the following variables: sbp (power = -2), tobacco (power = 0.4), ldl (power = 0.1), obesity (power = -0.4), and alcohol (power = 0.4)

```

[ ]: from scipy import stats
df = data[['sbp', 'tobacco', 'ldl', 'adiposity', 'typea', 'obesity', 'alcohol', 'age']]

lambda_values = [-2.0, 0.4, 0.1, 1.0, 1.0, -0.4, 0.4, 1.0]

df_power = df.copy()

for column, lam in zip(df, lambda_values):
    df_power[column] = stats.boxcox(df[column], lmbda=lam)

```

11 7. Produce histogram of each of the following transformed variables with imposing normal curve: sbp, tobacco, ldl, obesity, and alcohol.

```

[ ]: df_power_red = df_power[['sbp', 'tobacco', 'ldl', 'obesity', 'alcohol']]

num_cols = len(df_power_red.columns)

```

```

fig, axes = plt.subplots(3, 3, figsize=(10, 8))

axes = axes.ravel()

for i in range(min(num_cols, 9)):
    column_data = df_power_red.iloc[:, i]

    axes[i].hist(column_data, bins=10, density=True, alpha=0.6, color='b')

    mu, std = column_data.mean(), column_data.std()
    xmin, xmax = axes[i].get_xlim()
    x = np.linspace(xmin, xmax, 100)
    p = norm.pdf(x, mu, std)

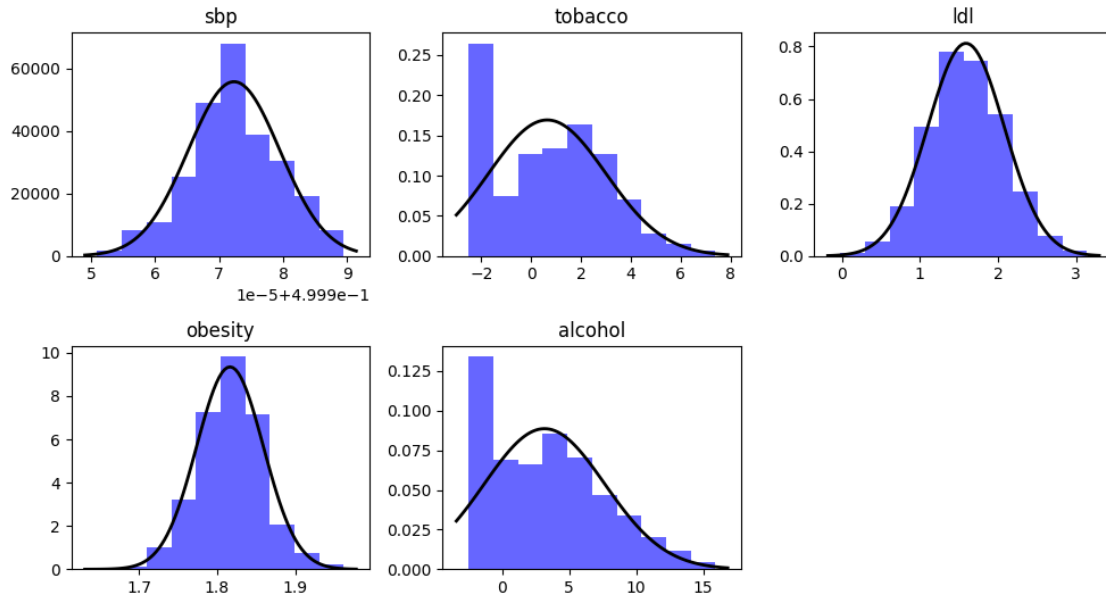
    axes[i].plot(x, p, 'k', linewidth=2)

    axes[i].set_title(df_power_red.columns[i])

for i in range(num_cols, 9):
    axes[i].axis('off')

plt.tight_layout()
plt.show()

```

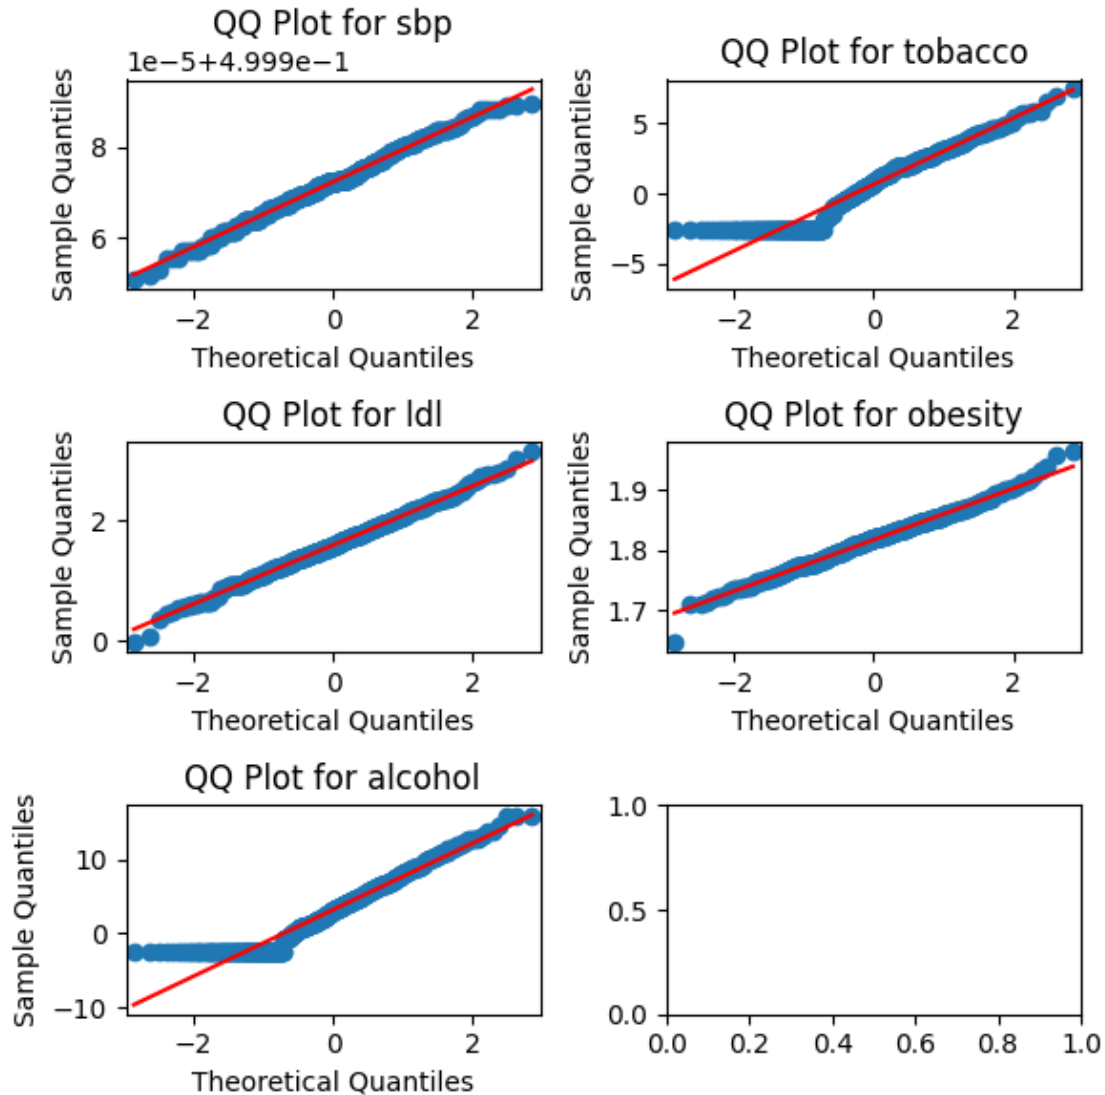


12 8. Produce quantile plot of each of the following transformed variables: sbp, tobacco, ldl, obesity, and alcohol

```
[ ]: fig, axes = plt.subplots(3, 2, figsize=(6, 6))
axes = axes.ravel()

for i, column in enumerate(df_power_red.columns):
    ax = axes[i]
    sm.qqplot(df_power_red[column], line='s', ax=ax)
    ax.set_title(f'QQ Plot for {column}')

plt.tight_layout()
plt.show()
```

13 9. Build a logistic regression model with all predictors (transformed and three remaining original variables do not perform any transformation)

```
[ ]: X_1 = df_power
      y_1 = data['chd']

      X_train_1, X_test_1, y_train_1, y_test_1 = train_test_split(X_1, y_1,
      ↪test_size=0.2, random_state=42)
```

```
# in LogisticRegression x=10.0 changes the regularization strength
model_1 = LogisticRegression(solver='liblinear', random_state=0).fit(X_1,y_1)
model_1.fit(X_train_1, y_train_1)

y_pred_1 = model_1.predict(X_test_1)

confusion_matrix(y_test_1, y_pred_1)
```

```
[ ]: array([[48, 11],
           [20, 14]])
```

```
[ ]: print(classification_report(y_test_1, y_pred_1))
```

	precision	recall	f1-score	support
0	0.71	0.81	0.76	59
1	0.56	0.41	0.47	34
accuracy			0.67	93
macro avg	0.63	0.61	0.62	93
weighted avg	0.65	0.67	0.65	93

14 10. Build another logistic regression model with all predictors as in Part 9 except using significant predictors only.

```
[ ]: X_2 = df_power_red
      y_2 = data['chd']

      X_train_2, X_test_2, y_train_2, y_test_2 = train_test_split(X_2, y_2,
      ↪test_size=0.2, random_state=42)

      # in LogisticRegression x=10.0 changes the regularization strength
      model_2 = LogisticRegression(solver='liblinear', random_state=0).fit(X_2,y_2)
      model_2.fit(X_train_2, y_train_2)

      y_pred_2 = model_2.predict(X_test_2)

      confusion_matrix(y_test_2, y_pred_2)
```

```
[ ]: array([[50, 9],
           [21, 13]])
```

```
[ ]: print(classification_report(y_test_2, y_pred_2))
```

	precision	recall	f1-score	support
0	0.70	0.85	0.77	59
1	0.59	0.38	0.46	34
accuracy			0.68	93
macro avg	0.65	0.61	0.62	93
weighted avg	0.66	0.68	0.66	93

PART II Reporting (10 Points) 1. After completion of this activity, complete the following table

```
[ ]: df_select_columns = data[['names', 'sbp', 'tobacco', 'ldl', 'adiposity',
    ↳ 'typea', 'obesity', 'alcohol', 'age']]

column_means = df_select_columns.mean()
column_median = df_select_columns.median()
column_skewness = df_select_columns.skew()

data_tbl = {
    'Variables': ['names', 'sbp', 'tobacco', 'ldl', 'adiposity', 'typea',
    ↳ 'obesity', 'alcohol', 'age'],
    'Mean':
    ↳ [column_means['names'], column_means['sbp'], column_means['tobacco'], column_means['ldl'], column_means['adiposity'],
    ↳ column_means['typea'], column_means['obesity'], column_means['alcohol'], column_means['age']],
    'Median':
    ↳ [column_median['names'], column_median['sbp'], column_median['tobacco'], column_median['ldl'], column_median['adiposity'],
    ↳ column_median['typea'], column_median['obesity'], column_median['alcohol'], column_median['age']],
    'Skewness':
    ↳ [column_skewness['names'], column_skewness['sbp'], column_skewness['tobacco'], column_skewness['ldl'], column_skewness['adiposity'],
    ↳ column_skewness['typea'], column_skewness['obesity'], column_skewness['alcohol'], column_skewness['age']]
}

Table = pd.DataFrame(data_tbl)

style_dict = {
    'text-align': 'center',
    'border': '1px solid black'
}

styled_Table = Table.style.set_properties(**style_dict).
    ↳ set_table_styles([{'selector': '', 'props': [('border', '1px solid
    ↳ black')]}])

styled_Table
```

```
[ ]: <pandas.io.formats.style.Styler at 0x7dfa7045a560>
```

2. Display the histogram and quantile plot of “tobacco”.

```
[ ]: tobacco_data = df_select_columns["tobacco"]

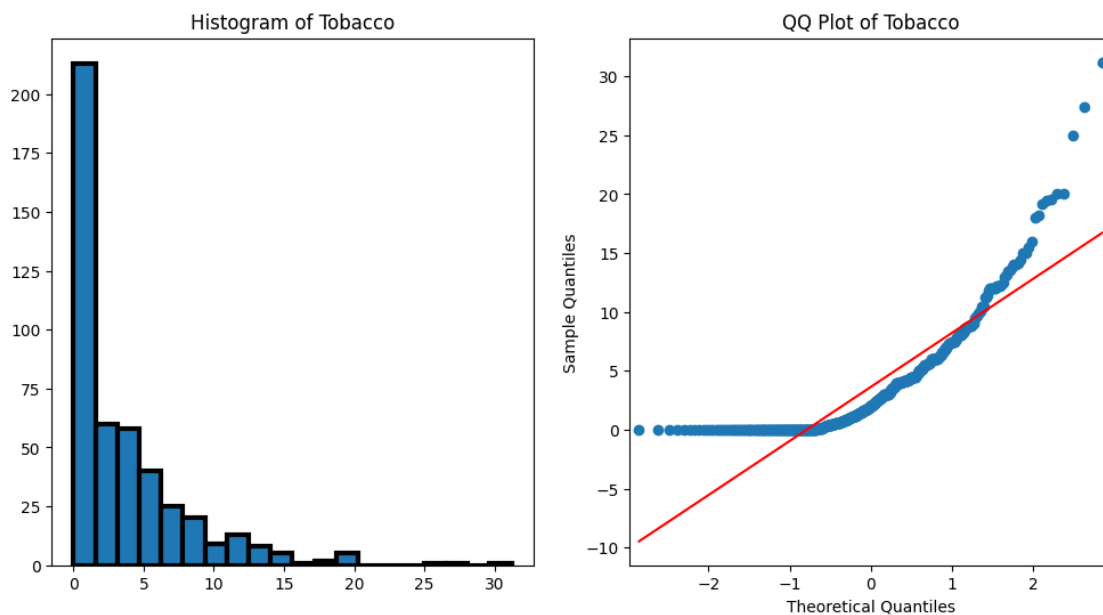
fig, axes = plt.subplots(1, 2, figsize=(12, 6))

axes[0].hist(tobacco_data, bins=20, edgecolor="k",linewidth=3)

axes[0].set_title("Histogram of Tobacco")

sm.qqplot(tobacco_data, line='s', ax=axes[1])
axes[1].set_title("QQ Plot of Tobacco")

plt.show()
```



3. Display the histogram and quantile plot of “alcohol”

```
[ ]: alcohol_data = df_select_columns["alcohol"]

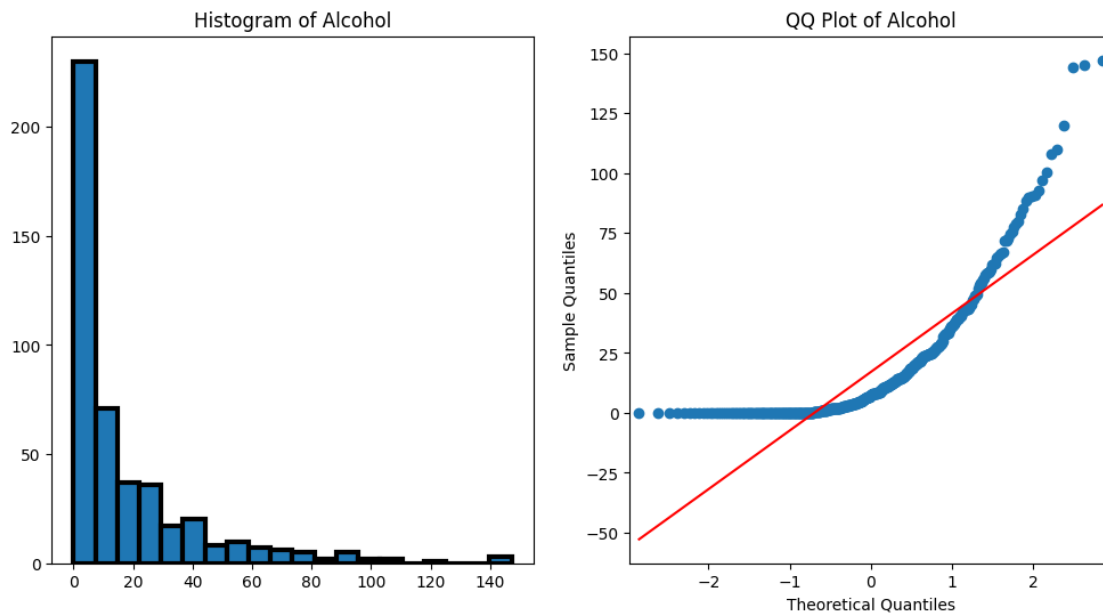
fig, axes = plt.subplots(1, 2, figsize=(12, 6))

axes[0].hist(alcohol_data, bins=20, edgecolor="k",linewidth=3)

axes[0].set_title("Histogram of Alcohol")
```

```
sm.qqplot(beer_data, line='s', ax=axes[1])
axes[1].set_title("QQ Plot of Beer")

plt.show()
```



4. Complete the following table

```
[ ]: for column in df_power.columns:
    data = df_power[column]
    mean = np.mean(data)
    median = np.median(data)
    skewness = stats.skew(data)

    print(f"Column: {column}")
    print(f"Mean: {mean}")
    print(f"Median: {median}")
    print(f"Skewness: {skewness}")
    print()
```

Column: sbp

Mean: 0.4999723389752801

Median: 0.4999721541546001

Skewness: -0.06309804169275851

Column: tobacco

Mean: 0.6524242122443431

Median: 0.7987697769322355
Skewness: 0.13429048073980526

Column: ldl
Mean: 1.5908671072190528
Median: 1.5810776346869808
Skewness: 0.02865204784191533

Column: adiposity
Mean: 24.4067316017316
Median: 25.115000000000002
Skewness: -0.21394839672197175

Column: typea
Mean: 52.103896103896105
Median: 52.000000000000001
Skewness: -0.34531194082632766

Column: obesity
Mean: 1.8166002789842062
Median: 1.8188253647320982
Skewness: 0.0002722556150250647

Column: alcohol
Mean: 3.158923226305479
Median: 3.099983437191029
Skewness: 0.39964197601136525

Column: age
Mean: 41.816017316017316
Median: 43.999999999999999
Skewness: -0.3804937421038407

7. 95% confidence interval for tobacco

```
[ ]: tobacco_data = df['tobacco']  
  
# Calculate the confidence interval  
results = sm.stats.DescrStatsW(tobacco_data)  
confidence_interval = results.tconfint_mean(alpha=0.05)  
  
print("95% Confidence Interval for 'tobacco':", confidence_interval)
```

95% Confidence Interval for 'tobacco': (3.215728421332208, 4.055570279966494)

8. 95% confidence interval for alcohol

```
[ ]: alcohol_data = df['alcohol']

# Calculate the confidence interval
results = sm.stats.DescrStatsW(alcohol_data)
confidence_interval = results.tconfint_mean(alpha=0.05)

print("95% Confidence Interval for 'tobacco':", confidence_interval)
```

95% Confidence Interval for 'tobacco': (14.806193411597095, 19.28259446719079)

9. Which model performed better based on C-statistic

Model 1

```
[ ]: from sklearn.metrics import roc_auc_score, roc_curve, auc
logreg = model
logreg.fit(X_train, y_train)

# Predict probabilities for the positive class (class 1) on the test set
y_pred_proba = logreg.predict_proba(X_test_1)[:, 1]

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

# Calculate the AUC (Area Under the Curve)
roc_auc = auc(fpr, tpr)

print(f'AUC-ROC (C-Statistic): {roc_auc:.2f}')
```

AUC-ROC (C-Statistic): 0.73

Model 2

```
[ ]: from sklearn.metrics import roc_auc_score, roc_curve, auc
logreg_1 = model_1
logreg_1.fit(X_train_1, y_train_1)

# Predict probabilities for the positive class (class 1) on the test set
y_pred_proba_1 = logreg_1.predict_proba(X_test_1)[:, 1]

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test_1, y_pred_proba_1)

# Calculate the AUC (Area Under the Curve)
roc_auc = auc(fpr, tpr)

print(f'AUC-ROC (C-Statistic): {roc_auc:.2f}')
```

AUC-ROC (C-Statistic): 0.78

Based on the C-statistic model 2 is the better model